

Problem Set #2

Experiment Design: Alex & Coco

May 22, 2016

Some guidelines for submitting problem sets in this course:

- Please submit a PDF document rather than a Word document or a Google document.
- Please put your name at the top of your problem set.
- Please **bold** or *highlight* your numerical answers to make them easier to find.
- If you'll be using **R** or **Python** code to calculate your answers, please put the code and its output directly into your Problem Set PDF document.
- It is highly recommended, although not required, that you use the RMarkdown feature in RStudio to compose your problem set answers. RMarkdown allows you to easily intermingle analysis code and answers in one document. It is of a similar design as **jupyter** and an **ipython** notebook.
- You do not need to show work for trivial calculations, but showing work is always allowed.
- For answers that involve a narrative response, please feel free to describe the key concept directly and briefly, if you can do so, and do not feel pressure to go on at length.
- Please ask us questions about the problem set if you get stuck. **Don't spend more than 20 minutes puzzling over what a problem means.**
- Please ensure that someone (us!) can compile your solution set. The best way is to use the web-hosted links we've provided.

FE exercise 3.6

The Clingingsmith, Khwaja, and Kremer study discussed in section 3.5 may be used to test the sharp null hypothesis that winning the visa lottery for the pilgrimage to Mecca had no effect on the views of Pakistani Muslims toward people from other countries. Assume that the Pakistani authorities assigned visas using complete random assignment.

```
## success views_saudi views_indonesian views_turkish views_african
## 1      0          1          1          0          0
## 2      0          1          1          0         -1
## 3      0          0          0          0          0
## 4      0          2          2          0          0
## 5      0          1          1          1          1
## 6      0          2          0          0          0
## views_chinese views_european views
## 1            0            0    2
## 2            1           -1    1
## 3            0            0    0
## 4            1            0    5
## 5            1           -2    3
## 6            0            0    2
```

- Conduct 10,000 simulated random assignments under the sharp null hypothesis.
- How many of the simulated random assignments generate an estimated ATE that is at least as large as the actual estimate of the ATE?

- c. What is the implied one-tailed p-value?
- d. How many of the simulated random assignments generate an estimated ATE that is at least as large *in absolute value* as the actual estimate of the ATE?
- e. What is the implied two-tailed p-value?

FE exercise 3.8

Naturally occurring experiments sometimes involve what is, in effect, block random assignment. For example, Titunik studies the effect of lotteries that determine whether state senators in TX and AR serve two-year or four-year terms in the aftermath of decennial redistricting. These lotteries are conducted within each state, and so there are effectively two distinct experiments on the effects of term length. An interesting outcome variable is the number of bills (legislative proposals) that each senator introduces during a legislative session. The table below lists the number of bills introduced by senators in both states during 2003.

```
library(foreign)
d3.8 <- read.dta("http://hdl.handle.net/10079/s1rn910")
head(d3.8)
```

##	term2year	bills_introduced	texas0_arkansas1
## 1	0	18	0
## 2	0	29	0
## 3	0	41	0
## 4	0	53	0
## 5	0	60	0
## 6	0	67	0

- a. For each state, estimate the effect of having a two-year term on the number of bills introduced.
- b. For each state, estimate the standard error of the estimated ATE.
- c. Use equation (3.10) to estimate the overall ATE for both states combined.
- d. Explain why, in this study, simply pooling the data for the two states and comparing the average number of bills introduced by two-year senators to the average number of bills introduced by four-year senators leads to biased estimate of the overall ATE.
- e. Insert the estimated standard errors into equation (3.12) to estimate the stand error for the overall ATE.
- f. Use randomization inference to test the sharp null hypothesis that the treatment effect is zero for senators in both states.
- g. **IN Addition:** Plot histograms for both the treatment and control groups in each state (for 4 histograms in total).

FE exercise 3.11

Use the data in table 3.3 to simulate cluster randomized assignment. (*Notes: (a) Assume 3 clusters in treatment and 4 in control; and (b) When Gerber and Green say “simulate”, they do not mean “run simulations with R code”, but rather, in a casual sense “take a look at what happens if you do this this way.” There is no randomization inference necessary to complete this problem.*)

- Suppose the clusters are formed by grouping observations $\{1,2\}$, $\{3,4\}$, $\{5,6\}$, \dots , $\{13,14\}$. Use equation (3.22) to calculate the standard error assuming half of the clusters are randomly assigned to treatment.
- Suppose that clusters are instead formed by grouping observations $\{1,14\}$, $\{2,13\}$, $\{3,12\}$, \dots , $\{7,8\}$. Use equation (3.22) to calculate the standard error assuming half of the clusters are randomly assigned to treatment.
- Why do the two methods of forming clusters lead to different standard errors? What are the implications for the design of cluster randomized experiments?

More Practice #1

You are an employee of a newspaper and are planning an experiment to demonstrate to Apple that online advertising on your website causes people to buy iPhones. Each site visitor shown the ad campaign is exposed to \$0.10 worth of advertising for iPhones. (Assume all users could see ads.) There are 1,000,000 users available to be shown ads on your newspaper's website during the one week campaign.

Apple indicates that they make a profit of \$100 every time an iPhone sells and that 0.5% of visitors to your newspaper's website buy an iPhone in a given week in general, in the absence of any advertising.

- By how much does the ad campaign need to increase the probability of purchase in order to be “worth it” and a positive ROI (supposing there are no long-run effects and all the effects are measured within that week)?
 - Assume the measured effect is 0.2 percentage points. If users are split 50:50 between the treatment group (exposed to iPhone ads) and control group (exposed to unrelated advertising or nothing; something you can assume has no effect), what will be the confidence interval of your estimate on whether people purchase the phone?
- Note:** The standard error for a two-sample proportion test is $\sqrt{p(1-p) * (\frac{1}{n_1} + \frac{1}{n_2})}$ where $p = \frac{x_1 + x_2}{n_1 + n_2}$, where x and n refer to the number of “successes” (here, purchases) over the number of “trials” (here, site visits). The length of each tail of a 95% confidence interval is calculated by multiplying the standard error by 1.96.
- Is this confidence interval precise enough that you would recommend running this experiment? Why or why not?
 - Your boss at the newspaper, worried about potential loss of revenue, says he is not willing to hold back a control group any larger than 1% of users. What would be the width of the confidence interval for this experiment if only 1% of users were placed in the control group?

More Practice #2

Here you will find a set of data from an auction experiment by John List and David Lucking-Reiley (2000).

```
d2 <- read.csv("https://docs.google.com/spreadsheets/d/1M2wSeI4xB1YWoHAjzTNvmJHBLwbaewOf8FHMSWRQsb8/pub
head(d2)
```

##	bid	uniform_price_auction
## 1	5	1
## 2	5	1
## 3	20	0
## 4	0	1
## 5	20	1
## 6	0	1

In this experiment, the experimenters invited consumers at a sports card trading show to bid against one other bidder for a pair trading cards. We abstract from the multi-unit-auction details here, and simply state that the treatment auction format was theoretically predicted to produce lower bids than the control auction format. We provide you a relevant subset of data from the experiment.

- Compute a 95% confidence interval for the difference between the treatment mean and the control mean, using analytic formulas for a two-sample t-test from your earlier statistics course.
- In plain language, what does this confidence interval mean?
- Regression on a binary treatment variable turns out to give one the same answer as the standard analytic formula you just used. Demonstrate this by regressing the bid on a binary variable equal to 0 for the control auction and 1 for the treatment auction.
- Calculate the 95% confidence interval you get from the regression.
- On to p-values. What p-value does the regression report? Note: please use two-tailed tests for the entire problem.
- Now compute the same p-value using randomization inference.
- Compute the same p-value again using analytic formulas for a two-sample t-test from your earlier statistics course. (Also see part (a).)
- Compare the two p-values in parts (e) and (f). Are they much different? Why or why not? How might your answer to this question change if the sample size were different?